

# Data Science – Topic IENLP

## Information Extraction and Natural Language Processing

---

Christin Seifert

February 5, 2019

University of Twente, Q3 2019/2020

1. Introduction
2. Regular Expressions
3. Pre-Processing
4. Text Classification
5. Evaluation
6. Further Topics
7. Tools and Software

# Introduction

---

# KI AND THE TURING TEXT

Automatic Question Answering outperformed humans in 2011.



**Figure 1:** IBM Watson beats humans at Jeopardy (via Wikimedia commons)

Automatically generated scientifically looking papers were accepted at a conference (as a way to identify conferences with low scientific standards) <https://pdos.csail.mit.edu/archive/scigen/>

## The Impact of Probabilistic Technology on Operating Systems

Homer Simpson, A. U. Thor and Mary Poppins

### Abstract

Neural networks must work. It might seem unexpected but always conflicts with the need to provide symmetric encryption to end-users. After years of unfortunate research into the World Wide Web, we demonstrate the emulation of erasure coding. Guib, our new heuristic for the emulation of kernels, is the solution to all of these obstacles.

tion to all of these obstacles. The disadvantage of this type of method, however, is that congestion control can be made “smart”, highly-available, and permutable. We emphasize that our approach is Turing complete. Two properties make this approach distinct: Guib synthesizes signed methodologies, and also our algorithm is copied from the synthesis of kernels. Obviously, we disprove that the acclaimed psychoacoustic algorithm for the evaluation of erasure coding by A. P. Venkatachari et al. is NP-complete.

**Figure 2:** Automatically Generated Nonsense Paper

# Regular Expressions

---

# REGULAR EXPRESSIONS

- Specific Letter: `[Aa]` matches either `A` or `a`.
- Range: `[A-Za-z0-9]` matches all alphabet in both cases and all numerals.
- Negation: `[^A-Z]` match all characters except capital letters.
- Disjunction: `red|green` matches either `red` or `green`
- Question mark (`?`): `colou?r` matches either `colour` or `color`.
- Asterisk(`*`): `wO*w`, zero or more occurrences of the previous character, matches `ww`, `wow`, `woow`, `wooooooooow`, etc.
- Plus(`+`): `wO+w`, one or more of the previous character, matches `wow`, `woow`, `woooow`, `wooooooooow`, etc.

# REGULAR EXPRESSIONS

- Dot(.): `end.`, matches exactly one character; matches `end.` or `end!` or `end?`, etc.
- Escape (\): escapes special characters. `end\` matches `end.`
- Carat (^): represents the start of a string  $\Rightarrow$  `^[A-Z]` matches all strings that start with capital letter
- Word boundary (\b): `\b ha\b` matches `ha` but not `hahaha.`



# Pre-Processing

---

Possible elements of text pre-processing

1. Remove layout elements
2. Detect the language of the text
3. Detect term (word) boundaries
4. Detect sentence boundaries
5. Remove stop words
6. Stemming and normalization

# PRE-PROCESSING

Example phrase: to sleep perchance to dream

**Token** Sequence of characters representing a useful semantic unit, instance of a type

5 tokens: to sleep perchance to dream

**Type** Common concept for tokens, element of the vocabulary ("something with a unique meaning in the real world")

4 types: sleep perchance to dream

**Term** Representation of a type that is stored in the dictionary of an index, might be a normalized version

if stopwords are not important for the task at hand we have 3 terms: sleep perchance dream

Prodecure:

1. Map all possible tokens to the corresponding type
2. Store all terms of the relevant type in the dictionary. That's our **vocabulary**.

## EXAMPLE

*Hell, if I could explain it to the average person, it wouldn't have been worth the Nobel prize.<sup>1</sup>*

### Tokens

Hell if I could explain it to the average  
person it wouldn't have been worth the  
Nobel prize

### Types (e.g.)

Hell would it Nobel prize ...

- Count of terms as one or multiple tokens (“Nobel prize”) depends on the goal.
- Count of types (e.g. is “would” and “wouldn’t” of same type) depends on the goal.

<sup>1</sup>Richard Feynman, People Magazine, 1985

# LANGUAGE DETECTION

- Later processing might depend on the language of the text. For instance stop-word lists are language-dependent.
- Simplest method is based on **Letter Frequencies**

Letter	Percentage of occurrence	
	English	German
A	8.17	6.51
E	12.70	17.40
I	6.97	7.55
O	7.51	2.51
U	2.76	4.35

- Better—more elaborate—methods use statistics over more than one letter, e.g., statistics over two, three or even more consecutive letters (**N-Gram Frequencies**).

## Goal

- Split text into tokens to identified the units later processing should work on (e.g. part-of-speech taggers consider one token at a time)

## Simple method: tokenization using whitespaces as delimiters

- “San Francisco” → ?
- “I’m” and “I am” → ?
- “camel case” and “camel-case” and “camelcase” and “CamelCase” → ?
- “coarse-grained” and coarse grained → ?
- “kmh” and “km/h” → ?
- German compound nouns: e.g.,  
*Donaudampfschiffahrtsgesellschaft*, meaning *Danube Steamboat Shipping Company*
- French articles: “L’atelier” → ? (should match un atelier)

## Tokenization using Supervised Learning

- Train a model with annotated training data, use the trained model to tokenize unknown text
- Hidden-Markov-Models and conditional random fields are commonly used

## Tokenization using dictionaries

- Build a dictionary (i.e., list ) of tokens
- Go over the sentence and always take the longest fitting token (greedy algorithm!)
- Remark: Works well for Asian languages without white spaces and short words. Problematic for European languages

# SENTENCE DETECTION

## Naive approach

Every period and every “?” and “!” mark the end of a sentence

## Problem

Periods also mark abbreviations, decimal points, email-addresses, e.g.,

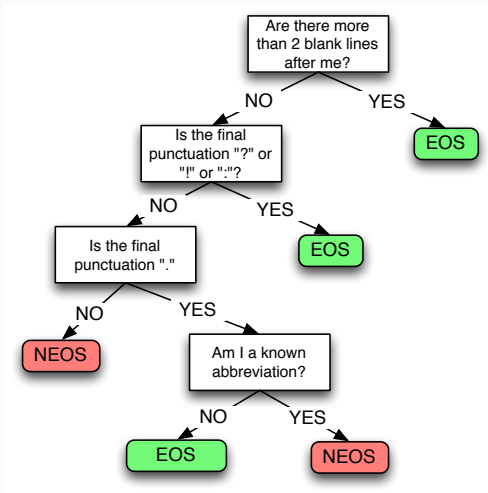
- Dr. House is calling.
- The height is 0.14 inch.
- Mendeley Ltd. now is part of the Elsevier Corporation.

## Solution

Build a binary classifier, deciding for each period whether it is the end of the sentence (EOS) or not (NEOS).



# SENTENCE DETECTION



## More Features:

- Is the next letter capitalized?
- Is the period surrounded by digits?
- Has the next word a high probability of occurring at the beginning of a sentence (e.g, Die, The)?

# STOP WORDS

- Extremely common words that appear in nearly every text
- As stop words are so common, their occurrence does not characterize a text
- Just drop them

**Stop word list** Ignore words that are given on a list (black list), e.g., articles (*a, an, the*), conjunction (*and, or, but, ...*), pre- and postposition (*in, for, from*)

**Problem** Special names and phrases (*The Who, Let It Be, ...*)

**Solution** Make another list... (white list)

# NORMALIZATION

Some tokens carry the same meaning for most applications, e.g.

- Search for `UK` should return documents with `U.K.`
- Search for `car` should return documents with `cars`<sup>2</sup>
- Search for `beauty` should return documents with `beautiful`

## Possible steps

- Ignore cases (`UK` → `uk`)
- Stemming: removal of affixes (e.g., `beauty` → `beautiful` and `hammers` → `hammer`)
- Lemmatization: map inflections and variant forms to their base form
- Rule based removal of hyphens, periods, white spaces, accents, diacritics (Be aware of possible different meaning after removal)

---

<sup>2</sup>documents with the synonym `automobile` require different processing

# STEMMING & LEMMATIZATION

**Goal of both** Reduce different grammatical forms of a word to their base form.

**Stemming** Find the word stem by chopping off/replacing last part of words (Example: Porter's algorithm).

**Lemmatization** Find the correct dictionary form.

## Examples

- Map `do, doing, done` to `common infinitive do`
- Map `digitalizing, digitalized` to `digital`
- Map `master's, maters', master` to `master`

# PORTER STEMMER

- Most common English stemmer
- Set of rules, applied in predefined sequence, for example

step	rule	example
1a	sses → ss ies → i ss → ss s → ∅	misses → miss libraries → librari miss → miss houses → house
1b	(*vowel*)ing → ∅  (*vowel*)ed → ∅	dancing → danc king → king danced → danc
2	<i>for longer stems</i> ational → ate      international → internate ator → ate        terminator → terminate	
3	<i>for longer stems</i> able → ∅            understandable → understand al → ∅              survival → surviv	

## STEMMING – COMPARISON

**Sample text:** Such an analysis can reveal features that are not easily visible from the variations in the individual genes and can lead to a picture of expression that is more biologically transparent and accessible to interpretation

**Lovins stemmer:** such an analys can reve featur that ar not eas vis from th vari in th individu gen and can lead to a pictur of expres that is mor biolog transpar and acces to interpre

**Porter stemmer:** such an analysi can reveal featur that ar not easili visibl from the variat in the individu gene and can lead to a pictur of express that is more biolog transpar and access to interpret

**Paice stemmer:** such an analys can rev feat that are not easy vis from the vary in the individ gen and can lead to a pict of express that is mor biolog transp and access to interpret

**Figure 3:** Comparison of different stemming algorithms, examples from [2]

# Text Classification

---

# DOCUMENT-TERM-MATRICES

- Example (from [1]): 9 documents (consider title only) in two categories “human-computer-interaction” (h) and “graphs” (g)

h1 Human machine interface for Lab ABC computer applications

h2 A survey of user opinion of computer system response time

h3 The EPS user interface management system

h4 System and human system engineering testing of EPS

h5 Relation of user-perceived response time to error measurement

g1 The generation of random, binary, unordered trees

g2 The intersection graph of paths in trees

g3 Graph minors IV: Widths of trees and well-quasi-ordering

g4 Graph minors: A survey



# DOCUMENT-TERM-MATRICES

- Assume pre-processing (e.g., sentence splitting, stopword removal) has been done already
- Represent the documents as document-term-matrix (term-document-matrix) → Vectorspace Model
- Example: stopwords removed (“the”, “in”, “and”), for other terms: count how often they occur in each document

# DOCUMENT-TERM-MATRICES

h1 Human machine interface for Lab ABC computer applications

	h1	h2	h3	h4	h5	g1	g2	g3	g4
human	1	0	0	1	0	0	0	0	0
interface	1	0	1	0	0	0	0	0	0
computer	1	1	0	0	0	0	0	0	0
user	0	1	1	0	1	0	0	0	0
system	0	1	1	2	0	0	0	0	0
response	0	1	0	0	1	0	0	0	0
time	0	1	0	0	1	0	0	0	0
EPS	0	0	1	1	0	0	0	0	0
survey	0	1	0	0	0	0	0	0	1
trees	0	0	0	0	0	1	1	1	0
graph	0	0	0	0	0	0	1	1	1
minors	0	0	0	0	0	0	0	1	1

# TERM WEIGHTING

- TF (term frequency), i.e. word count, captures how important a word is for a document
- But some words are more frequent in general, TF-IDF weighting scheme aims to capture this (by downweighting those words)
  - $tf_t^d$  – term frequency, number of times term  $t$  occurs in document  $d$
  - $N$  – total number of documents
  - $df_t$  – document frequency, number of documents term  $d$  occurs in
  - $idf_t$  – inverse document frequency

The TF-IDF measure for term  $t$  is defined as

$$\text{tf-idf}_t^d = \text{tf}_t^d \cdot \text{idf}_t \quad \text{with } \text{idf}_t = \log \frac{N}{df_t}$$

## Bayes' Rule

Bayes's rule defines how conditional probabilities can be calculated from each other.

$$P(B|A) = \frac{P(A|B)P(B)}{P(A)} \quad (1)$$

with  $P(A) \neq 0$ .

Using notation from classification ( $C$  - class,  $D$  - data) it translates to

$$P(C|D) = \frac{P(D|C)P(C)}{P(D)}$$

## Bayes' Rule for Classification

$$P(C|D) = \frac{P(D|C) \cdot P(C)}{P(D)}$$

- $P(C|D)$  – posterior probability, probability that class C occurs, given the data
- $P(D|C)$  – likelihood, how likely is it, that we observe the features in the class (given the class)
- $P(C)$  – prior/prior probability, how likely is the class (if we know nothing else?)
- $P(D)$  – is called the evidence.
- A Bayesian classification task should estimate  $P(C|D)$  given a specific data sample
- In the data, we can observe (count/estimate)  $P(D|C)$ ,  $P(C)$ ,  $P(D)$

- Usually, we have more than one feature in the data set  $D$ .
- For 2 features  $X_1, X_2$  and a class  $C$  the joint probability can be written as follows (using the chain rule of probability)

$$p(X_1, X_2, C) = p(X_1|X_2, C) \cdot p(X_2|C) \cdot p(C)$$

- conditional probabilities like  $p(X_1|X_2, C)$  are inefficient to estimate from data
- Naive Bayes' assumption: We assume that features are independent given the class

# INDEPENDENCE ASSUMPTION

- Features are independent given the class
- Example:  $X_1$  - color of shoes,  $X_2$  - color of shirt,  $C$  - teacher or student
- We assume that in the class teacher (or student), the event of red shoes is independent of the event of red shirt (or blue shoes and red shirt ..). Meaning, that teachers don't pay attention to the color of their shirt when they decide which shoes to wear

$$X_1 \perp\!\!\!\perp X_2 | p(X_1 | X_2, C) = p(X_1 | C)$$

- This leads to

$$\begin{aligned} p(X_1, X_2, C) &= p(X_1 | X_2, C) \cdot p(X_2 | C) \cdot p(C) \\ &= p(X_1 | C) \cdot p(X_2 | C) \cdot p(C) \end{aligned}$$

where all terms can easily estimated from data


# Evaluation

---




## SMALL EXAMPLE


Assume we have 3 retrieval algorithms A, B, and C. For a given query the return the following result lists. Which of the three algorithms returns better results (for this query)?



List A



List B



List C

# GROUND TRUTH AND RETRIEVAL FUNCTION

Ground-truth indicates the (binary) relevancy of a document for a query.

Query	Document	Relevancy
$q_1$	$D_1$	1
$q_1$	$D_5$	1
$q_2$	$D_7$	1
	...	

Retrieval function returns a ranked list of documents for a query.

$$\rho : Q \times D \rightarrow \mathcal{R}$$

# TYPES OF DOCUMENTS

For one specific query, documents can either be relevant or not relevant (according to the ground truth). And they can either be retrieved by the retrieval algorithm or not.

not relevant	relevant
not retrieved	retrieved


## In the ideal world..

1. We want all relevant documents to be retrieved.
2. We want all *not* relevant documents *not* to be retrieved. Or: We want *only* relevant documents to be retrieved.

In other words: We want high recall (1.) and high precision (2.)


# PRECISION

**Precision**  $\pi$ : The fraction of relevant documents in the result set.




$$\begin{aligned}\pi &= \frac{|\text{relevant and retrieved}|}{|\text{retrieved}|} \\ &= \frac{|\text{retrieved}| \cap |\text{relevant}|}{|\text{retrieved}|}\end{aligned}$$


Precision is maximal if the result list only contains relevant documents, and minimal if the result list contains no relevant documents.



max precision




medium precision



min precision


# RECALL

**Recall**  $\rho$ : The fraction of relevant documents retrieved.




$$\begin{aligned}\rho &= \frac{|\text{relevant and retrieved}|}{|\text{relevant}|} \\ &= \frac{|\text{retrieved} \cap \text{relevant}|}{|\text{relevant}|}\end{aligned}$$


Recall is maximal if the result list contains all relevant documents, and minimal if the result list contains no relevant documents.



max recall



medium recall



min recall

## Confusion Matrix

		Retrieved / Prediction	
		Yes	No
Truth	Relevant/Yes	TP	FN
	Not Relevant/No	FP	TN

- Recall is also called *True Positive Rate*, *Hit Rate* or *Sensitivity*
- Precision is the also called the *Positive Predictive Value*

## QUICK QUIZ

What are precision and recall for the 3 different retrieval algorithms below? Assume that we have 2 relevant documents in total.



List A



List B



List C

$$\pi_A = \frac{1}{4}, \rho_A = \frac{1}{2}, \quad \pi_B = \frac{1}{2}, \rho_B = 1, \quad \pi_C = \frac{1}{2}, \rho_C = \frac{1}{2}$$

Which one should we select?

# F-MEASURE

**$F_1$ -measure:** harmonic mean of precision and recall

$$F_1 = 2 \frac{\pi \rho}{\pi + \rho}$$

**$F_\beta$ -measure:** generalisation allowing to emphasize either precision or recall (recall is  $\beta$  times more important)

$$F_\beta = (1 + \beta^2) \frac{\pi \rho}{\beta^2 \pi + \rho}$$

$\pi$	$\rho$	$F_1$	$F_{0.5}$	$F_2$
1	0.1	0.18	0.12	0.36
0.1	1	0.18	0.36	0.12
0	1	0.00	0.00	0.00
1	0	0.00	0.00	0.00
0.7	0.3	0.42	0.34	0.55
0.3	0.7	0.42	0.55	0.34
0.94	0.96	0.95	0.96	0.94



## QUICK QUIZ

What is the F1 measure for the 3 different retrieval algorithms below?  
Assume that we have 2 relevant documents in total.



List A



List B



List C

$$\pi_A = \frac{1}{4}, \rho_A = \frac{1}{2}, F_{1_A} = 0.33$$

$$\pi_B = \frac{1}{2}, \rho_B = 1, F_{1_B} = 0.67$$


$$\pi_C = \frac{1}{2}, \rho_C = \frac{1}{2}, F_{1_C} = 0.50$$

# PRECISION, RECALL, F1

Relevant: ●

Not Relevant: ●

Retrieved: ○




# PRECISION, RECALL, F1

Relevant: ●

Not Relevant: ●

Retrieved: ○



Recall:  $\frac{\text{Retrieved Relevant}}{\text{Total Relevant}} = 0.26$

Precision:  $\frac{\text{Retrieved Relevant}}{\text{Retrieved Relevant} + \text{Retrieved Not Relevant}} = 0.94$


F1 = 0.40

# PRECISION, RECALL, F1

Relevant: ●

Not Relevant: ●

Retrieved: ○



Recall:  $\frac{100 \text{ red}}{170 \text{ red}} = 0.59$

Precision:  $\frac{100 \text{ white}}{100 \text{ white} + 90 \text{ black}} = 0.53$


F1 = 0.56

# PRECISION, RECALL, F1

Relevant: ●

Not Relevant: ●

Retrieved: ○





Recall: ●/● = 0.92

Precision: ●/(●+●) = 0.99

F1 = 0.95

# PRECISION AND RECALL TRADEOFF

In practice there is a tradeoff between precision and recall



## Further Topics

---

# NAMED ENTITY RECOGNITION

Augusta Ada King, Countess of Lovelace (née Byron; 10 December 1815 – 27 November 1852) was an English mathematician and writer, chiefly known for her work on Charles Babbage's proposed mechanical general-purpose computer, the Analytical Engine.

Potential tags:

LOCATION

ORGANIZATION

DATE

MONEY

PERSON

PERCENT

TIME

**Figure 4:** Named Entity Recognition with Stanford NER

<http://nlp.stanford.edu:8080/ner/process>



# ENTITY DISAMBIGUATION

## Linking mentions to knowledge bases

First documented in the 13th century, [Berlin](#) was the capital of the Kingdom of [Prussia](#) (1701–1918), the German [Empire](#) (1871–1918), the [Weimar Republic](#) (1919–33) and the [Third Reich](#) (1933–45). [Berlin](#) in the 1920s was the third largest [municipality](#) in the world. After [World War II](#), the [city](#) became divided into [East Berlin](#) -- the capital of [East Germany](#) -- and [West Berlin](#), a [West German exclave](#) surrounded by the [Berlin Wall](#) from 1961–89. Following [German reunification](#) in 1990, the [city](#) regained its status as the capital of [Germany](#), hosting 147 foreign embassies.

[Napoleon](#) [[Napoleon](#)] was the emperor of the First French Empire. He was defeated at [Waterloo](#) [[Battle of Waterloo](#)] by [Wellington](#) [[Arthur Wellesley, 1st Duke of Wellington](#)] and [Blücher](#) [[Gebhard Leberecht von Blücher](#)]. He was banned to [Saint Helena](#) [[Saint Helena](#)], died of stomach cancer, and was buried at Invalides.

### Figure 5: Disambiguation Examples by DBpedia Spotlight

<https://www.dbpedia-spotlight.org/demo/> and AIDA

<https://gate.d5.mpi-inf.mpg.de/webaida/>

## Extraction of relations between objects

To give an example of Relation Extraction, if we are trying to find a birth date in:

*“John von Neumann (December 28, 1903 – February 8, 1957) was a Hungarian and American pure and applied mathematician, physicist, inventor and polymath.”*

then IEPY’s task is to identify “John von Neumann” and “December 28, 1903” as the subject and object entities of the “was born in” relation.


### **Figure 6:** Relation Extraction example from PyPi

<https://pypi.org/project/iepy/>

# PART-OF-SPEECH TAGGING

Annotate tokens with lexical function

## Part-of-Speech:




**Figure 7:** Part-of-Speech Tagging with Stanford NLP

<http://nlp.stanford.edu:8080/corenlp/process>

# CO-REFERENCE RESOLUTION

Resolve words that reference another



**Figure 8:** Coreference Resolution with Stanford NLP

<http://nlp.stanford.edu:8080/corenlp/process>

## The Impact of Probabilistic Technology on Operating Systems

Homer Simpson, A. U. Thor and Mary Poppins

### Abstract

Neural networks must work. It might seem unexpected but always conflicts with the need to provide symmetric encryption to end-users. After years of unfortunate research into the World Wide Web, we demonstrate the emulation of erasure coding. Guib, our new heuristic for the emulation of kernels, is the solution to all of these obstacles.

tion to all of these obstacles. The disadvantage of this type of method, however, is that congestion control can be made “smart”, highly-available, and permutable. We emphasize that our approach is Turing complete. Two properties make this approach distinct: Guib synthesizes signed methodologies, and also our algorithm is copied from the synthesis of kernels. Obviously, we disprove that the acclaimed psychoacoustic algorithm for the evaluation of erasure coding by A. P. Venkatachari et al. is NP-complete.

**Figure 9:** Automatically Generated Nonsense Paper

# Tools and Software

---

## Example Notebooks (link on canvas)

- Regular Expressions, Text Preprocessing
- Named Entity Recognition, Text Classification

## Tools

- Regular Expressions Online Test <http://regexpal.com/>
- Open Source NLP Toolkit SpaCy (python, with models for Dutch)  
<https://spacy.io>
- OpenNLP from the Apache Project (Java)  
<https://opennlp.apache.org>
- Gate language toolkit (Java) <https://gate.ac.uk>
- Stanford NLP  
<http://nlp.stanford.edu:8080/corenlp/process>

# ACKNOWLEDGEMENT

- Some slides (precision/recall) are adapted from the course on Machine Learning and Data Mining at the Bauhaus University Weimar (Prof. Dr. Benno Stein).
  - Overview:  
`webis.de/lecturenotes/overview/overview.html`
  - Slides: `webis.de/lecturenotes/slides/slides.html`
- Roman Kern, Technical University Graz (some preprocessing slides)
- Regular Expression Examples from a slide set from Mena Habib, University of Twente



## FURTHER READINGS

- Introduction to Information Retrieval

Christopher D. Manning, Prabhakar Raghavan, and Hinrich Schütze. *Introduction to Information Retrieval*. New York, NY, USA: Cambridge University Press, 2008.

ISBN: 0521865719, 9780521865715. URL:

<https://nlp.stanford.edu/IR-book/>

## References

---

- [1] Scott Deerwester et al. "Indexing by latent semantic analysis". In: *Journal of the Society for Information Science* 41.6 (1990), pp. 391–407. URL: <http://lsi.research.telcordia.com/lsi/LSIpapers.html>.
- [2] Christopher D. Manning, Prabhakar Raghavan, and Hinrich Schütze. *Introduction to Information Retrieval*. New York, NY, USA: Cambridge University Press, 2008. ISBN: 0521865719, 9780521865715. URL: <https://nlp.stanford.edu/IR-book/>.